

```

1 from google.colab import files
2 import pandas as pd
3
4 # Upload the CSV file
5 uploaded = files.upload()

```



Choose Files spam.csv

- spam.csv(text/csv) - 503663 bytes, last modified: 9/20/2019 - 100% done
- Saving spam.csv to spam.csv

```

1 # Load the dataset
2 df = pd.read_csv('spam.csv', encoding='latin1')
3
4 # Display the first few rows of the dataframe
5 df.head()

```



	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

Next steps:

[Generate code with df](#)

[View recommended plots](#)

[New interactive sheet](#)

```

1 from sklearn.model_selection import train_test_split
2 from sklearn.feature_extraction.text import TfidfVectorizer

```

```
1 df.info()
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0    v1          5572 non-null   object
1    v2          5572 non-null   object
2    Unnamed: 2   50 non-null     object
3    Unnamed: 3   12 non-null     object
4    Unnamed: 4    6 non-null     object
dtypes: object(5)
memory usage: 217.8+ KB

```

```

1 # Drop unnecessary columns
2 df = df[['v1', 'v2']]
3
4 # Rename columns for convenience
5 df.columns = ['label', 'message']
6
7 # Encode the labels: 'spam' -> 1, 'ham' -> 0
8 df['label'] = df['label'].map({'ham': 0, 'spam': 1})

```



```

<ipython-input-12-f40901fbba71>:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-vs-returning-a-copy

```
df['label'] = df['label'].map({'ham': 0, 'spam': 1})
```

```

1 # Check the cleaned data
2 print(df.head())

```



	label	message
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```

1 # Split the data into features and target
2 X = df['message']
3 y = df['label']
4
5 # Split the data into training and test sets
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

```

```

1 # Initialize the TF-IDF Vectorizer
2 tfidf_vectorizer = TfidfVectorizer(stop_words='english')
3
4 # Fit and transform the training data
5 X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
6
7 # Transform the test data
8 X_test_tfidf = tfidf_vectorizer.transform(X_test)

```

```

1 from sklearn.naive_bayes import MultinomialNB
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.svm import SVC
4 from sklearn.metrics import classification_report

```

```

1 # Naive Bayes
2 nb_classifier = MultinomialNB()
3 nb_classifier.fit(X_train_tfidf, y_train)
4 y_pred_nb = nb_classifier.predict(X_test_tfidf)
5 print("Naive Bayes Classification Report:")
6 print(classification_report(y_test, y_pred_nb))

```

```

Naive Bayes Classification Report:

```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	1453
1	1.00	0.75	0.86	219
accuracy			0.97	1672
macro avg	0.98	0.88	0.92	1672
weighted avg	0.97	0.97	0.97	1672

```

1 # Logistic Regression
2 lr_classifier = LogisticRegression(max_iter=1000)
3 lr_classifier.fit(X_train_tfidf, y_train)
4 y_pred_lr = lr_classifier.predict(X_test_tfidf)
5 print("Logistic Regression Classification Report:")
6 print(classification_report(y_test, y_pred_lr))

```

```

Logistic Regression Classification Report:

```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	1453
1	0.98	0.65	0.78	219
accuracy			0.95	1672
macro avg	0.96	0.83	0.88	1672
weighted avg	0.95	0.95	0.95	1672

```

1 # Support Vector Machines (SVM)
2 svm_classifier = SVC()
3 svm_classifier.fit(X_train_tfidf, y_train)
4 y_pred_svm = svm_classifier.predict(X_test_tfidf)
5 print("SVM Classification Report:")
6 print(classification_report(y_test, y_pred_svm))

```




```

SVM Classification Report:

```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	1453
1	0.99	0.84	0.91	219
accuracy			0.98	1672
macro avg	0.99	0.92	0.95	1672
weighted avg	0.98	0.98	0.98	1672

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 from sklearn.metrics import confusion_matrix
4
5 # Function to plot confusion matrix
6 def plot_confusion_matrix(y_true, y_pred, title='Confusion Matrix'):
7     cm = confusion_matrix(y_true, y_pred)
8     plt.figure(figsize=(8, 6))
9     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
10                xticklabels=['Legitimate', 'Spam'], yticklabels=['Legitimate', 'Spam'])
11     plt.xlabel('Predicted')
12     plt.ylabel('True')
13     plt.title(title)
14     plt.show()
15
16 # Plot confusion matrix for Naive Bayes
17 plot_confusion_matrix(y_test, y_pred_nb, title='Confusion Matrix for Naive Bayes')
18
19 # Plot confusion matrix for Logistic Regression
20 plot_confusion_matrix(y_test, y_pred_lr, title='Confusion Matrix for Logistic Regression')
21
22 # Plot confusion matrix for SVM
23 plot_confusion_matrix(y_test, y_pred_svm, title='Confusion Matrix for SVM')
```



Legitimate

Predicted

Spam

```
1 from sklearn.metrics import roc_curve, auc
2
3 # Function to plot ROC curve
4 def plot_roc_curve(y_true, y_score, title='ROC Curve'):
5     fpr, tpr, _ = roc_curve(y_true, y_score)
6     roc_auc = auc(fpr, tpr)
7     plt.figure(figsize=(8, 6))
8     plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
9     plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
10    plt.xlim([0.0, 1.0])
11    plt.ylim([0.0, 1.05])
12    plt.xlabel('False Positive Rate')
13    plt.ylabel('True Positive Rate')
14    plt.title(title)
15    plt.legend(loc='lower right')
16    plt.show()
17
18 # Calculate probabilities for ROC curve
19 y_score_nb = nb_classifier.predict_proba(X_test_tfidf)[: , 1]
20 y_score_lr = lr_classifier.predict_proba(X_test_tfidf)[: , 1]
21 y_score_svm = svm_classifier.decision_function(X_test_tfidf)
22
23 # Plot ROC curves
24 plot_roc_curve(y_test, y_score_nb, title='ROC Curve for Naive Bayes')
25 plot_roc_curve(y_test, y_score_lr, title='ROC Curve for Logistic Regression')
26 plot_roc_curve(y_test, y_score_svm, title='ROC Curve for SVM')
```